

Lecture 14: Move Semantics

CS 106L, Fall '20

Last time...

- **Special Member Functions** get called for specific tasks
 - **Copy constructor:** create a new object as a **copy** of an existing object
`Type::Type(const Type& other)`
 - **Copy assignment:** reassign a new object to be a **copy** of an existing object
`Type::operator=(const Type& other)`
 - **Destructor:** deallocate the memory of an existing object
`Type::~Type()`
- SMFs are **automatically** generated for you
 - But if you're managing pointers to allocated to memory, do it yourself

Custom copy operator

We have to manually copy the elements themselves, not the pointer to array

```
template <typename T>
vector::vector<T>(const vector<T>& other) :
    _size(other._size),                // copy
    _capacity(other._capacity),        // copy
    _elems(new T[other._capacity]) {   // newly allocated
    std::copy(other._elems,
               other._elems + other._size, _elems);
}
```

Custom copy assignment

Fill in the blanks! What do we need to copy?

```
template <typename T>
vector& vector::operator=(const vector<T>& other) {
    // fill in the blanks
    // we need to copy other._size, other._capacity, and other._elems
    return *this;
}
```

Custom copy assignment

Fill in the blanks! What do we need to copy?

```
template <typename T>
vector& vector::operator=(const vector<T>& other) {
    _size = other._size;           // we can copy a size_t
    _capacity = other._capacity;  // we can copy a size_t
    resize(other._size);          // make sure we have enough space
    std::copy(other._elems,
              other._elems + other._size, _elems);
    return *this;
}
```

Answer in the chat

Which special member functions get called at each stage?

```
vector<string> findAllWords(int i);  
int main() {  
    vector<string> words1 = findAllWords(12345);    // (1)  
  
    vector<string> words2;                          // (2)  
  
    words2 = findAllWords(23456);                  // (3)  
}
```

Answer in the chat

Which special member functions get called at each stage?

```
int main() {  
    vector<string> words1 = findAllWords(12345); // (1)  
  
    vector<string> words2;  
    words2 = findAllWords(23456); // (2)  
}  
// (3)
```

copy constructor (points to the assignment operator in line 1)

destructor (points to the first `findAllWords` call in line 1)

copy assignment (points to the assignment operator in line 3)

destructor (points to the second `findAllWords` call in line 3)

Answer in the chat

Copy elision allows us to optimize out the first set of calls...

```
int main() {  
    vector<string> words1 = findAllWords(12345); // (1)  
  
    vector<string> words2;  
    words2 = findAllWords(23456); // (3)  
}
```

Annotations in the code:

- copy constructor** (red text) with an arrow pointing to the `=` operator on line 1.
- destructor** (green text) with an arrow pointing to the `findAllWords` call on line 1.
- copy assignment** (blue text) with an arrow pointing to the `=` operator on line 3.
- destructor** (green text) with an arrow pointing to the `findAllWords` call on line 3.



but we still can't optimize out making two **vectors** on line 3.

for now, let's ignore the effects of copy elision


The central problem

```
words2 = findAllWords(23456);
```

We need some way to **move** the resources of the result of **findAllWords** to **words2**, so that we don't create two objects (and immediately destroy one).


(Question in the chat: why don't we return **vector&** instead of **vector**?)

Copy



Why are we copying something we no longer need?

Moves



This works because we **no longer need** the temp value

One idea

What if **=** always represented a move? Does this work? (answer in chat)

```
int main() {  
    vector<string> words1 = findAllWords(12345);  
  
    vector<string> words2;  
    words2 = findAllWords(23456);  
}
```

Annotations:
- **move constructor** (points to the first **=**)
- **move assignment** (points to the second **=**)

✓ **Yes, this works correctly!** It is also efficient (no wasted memory).

Another example

What if **=** always represented a move? Does this work? (answer in chat)

```
int main() {  
    vector<string> words1 = findAllWords(12345);  
    vector<string> words2 = words1;  
  
    words1.push_back("Everything is fine!");  
}
```

Annotations:
- **move constructor** points to the first **=**
- **move assignment** points to the second **=**

 **No, this doesn't work!** words2 stole its array from words1 :(

Another example

To get this to work correctly, we need to use a combination...

```
int main() {  
    move constructor  
    vector<string> words1 = findAllWords(12345);  
    copy assignment  
    vector<string> words2 = words1;  
  
    words1.push_back("Everything is fine!");  
}
```

How can we distinguish between these two cases?

Topics

- lvalues vs. rvalues
- Move implementation
- Forcing a move to occur
- swap and insert

lvalues vs. rvalues

Looking more closely...

Recall: we could use **move** on the first line because it was **temporary**.

```
int main() {  
    vector<string> words1 = findAllWords(12345);  
    vector<string> words2 = words1;  
}
```

move constructor (red text) with an arrow pointing to the `=` operator in the first line.
copy assignment (blue text) with an arrow pointing to the `=` operator in the second line.

L-values and r-values generalize the idea of “temporariness.”

A **r-value** is “temporary,” and a **l-value** is not.

Official definition: a l-value has an address (can do &), and a r-value does not

A l-value can appear **left** or **right** of =.

```
x = 5;
```

```
y = x;
```

An r-value can only appear **right** of =.

```
5 = x;
```

```
y = 5;
```

Which of these are r-values? (Answer in chat)

```
int val = 2; // A
int* ptr = 0x02248837; // B
vector<int> v1{1, 2, 3}; // C
auto v4 = v1 + v2; // D
auto v5 = v1 += v4; // E
size_t size = v.size(); // F
val = static_cast<int>(size); // G
v1[1] = 4*i; // H
ptr = &val; // I
v1[2] = *ptr; // J
```

Which of these are r-values? (Answer in chat)

```
int val = 2; // A
int* ptr = 0x02248837; // B
vector<int> v1{1, 2, 3}; // C
auto v4 = v1 + v2; // D
auto v5 = v1 += v4; // E
size_t size = v.size(); // F
val = static_cast<int>(size); // G
v1[1] = 4*i; // H
ptr = &val; // I
v1[2] = *ptr; // J
```

A l-value's lifetime is until end of scope...

A r-value's lifetime is until end of line

unless you artificially extend it

References

A reference is an alias to an existing object

```
int main() {  
    vector<int> vec;  
    change(vec);  
}  
  
void change(vector<int>& v) {...}    // v is an alias of vec
```

Notice: vec is a **l-value**.

What happens when you try to reference a **r-value**?

```
int main() {  
    change(7);  
}
```

```
void change(int& a) {...}    // this doesn't work
```

- One caveat: in order to pass a variable by **reference**, you need to actually have a variable. The following does not work, for our example above:

```
doubleValueWithRef(15); // error! cannot pass a literal value by reference
```

Compiler error:

```
../all-examples.cpp:135:5: error: no matching function for call to 'doubleValueWithRef'
  doubleValueWithRef(15);
  ^~~~~~
../all-examples.cpp:11:6: note: candidate function not viable: expects an l-value for 1st argument
void doubleValueWithRef(int &x);
    ^
1 error generated.
```

l-value reference

```
int main() {  
    int i = 7;  
    change(i);  
}  
  
void change(int& a) {...}
```

r-value reference

```
int main() {  
    change(7);  
}  
  
void change(int&& a) {...}
```

Why rvalue references?

Copy constructor (lvalue& reference)

```
vector<T>(const vector<T>& other) :  
    _size(other._size),  
    _capacity(other._capacity) {  
    // have to manually copy the array  
    _elems = new T[other._capacity];  
    std::copy(other._elems,  
              other._elems + other._size,  
              _elems);  
}
```

Have to leave **other** in an usable state (**const**).

Move constructor (rvalue&& reference)

```
vector<T>(vector<T>&& other) :  
    _size(other._size),  
    _capacity(other._capacity) {  
    // steal the other array 🕵️  
    _elems = other._elems;  
  
    other._elems = nullptr;  
    other._size = 0;  
}
```

Brutally rob **other**'s resources, because it is temporary and will disappear when the function exits.

**Overloading between `&` and `&&` versions
of the same function allows us to
disambiguate between `copy` and `move`!**

but wait

An issue

Move constructor (rvalue&& reference)

```
vector<T>(vector<T>&& other) :  
    _size(other._size),  
    _capacity(other._capacity) {  
    // steal the other array 🕵️  
    _elems = other._elems;  
  
    other._elems = nullptr;  
    other._size = 0;  
}
```

this is a lvalue

(because **other&&**, a reference, is a lvalue)

so this performs... [answer in the chat]

a copy :(


Copying a pointer is cheap, but you could
imagine situations where this is a
problem

Fixing this

Move constructor (rvalue&& reference)

```
vector<T>(vector<T>&& other) :  
    _size(std::move(other._size)),  
    _capacity(std::move(other._capacity)) {  
    // steal the other array 🕵️  
    _elems = std::move(other._elems);  
  
    other._elems = nullptr;  
    other._size = 0;  
}
```

std::move is a cast to a rvalue&&
(equivalent to `std::static_cast<T&&>`)



Move Assignment

Move assignment (**rvalue&&** reference)

```
vector<T>& operator=(vector<T>&& other) {  
    _size = std::move(other._size);  
    _elems = std::move(other._elems);  
    _capacity = std::move(other._capacity);  
    other._elems = nullptr;  
    other._size = 0;  
}
```

```
vector<T>& v = {1, 2, 3};  
vector<T>& v2;  
v2 = std::move(v);
```

std::move is a cast to a rvalue&&
(equivalent to `std::static_cast<T&&>`)

Takeaways

- Use a **constructor** taking a **rvalue** for move constructor
- Use **operator=** taking a **rvalue** for move assignment
- Use **std::move** to make sure other object's values are treated as rvalues (and so moved)
 - Call **std::move** to force anything to become a rvalue (and get its data taken!)

std::move does not move anything

Recall...

A l-value's lifetime is until end of scope...

A r-value's lifetime is until end of line

unless you artificially extend it

For our vector...

```
template <typename T>
void vector<T>::push_back(const T& element) {
    elems[_size++] = element;           // equals → copy
}
```

```
template <typename T>
void vector<T>::push_back(T&& element) {
    elems[_size++] = std::move(element); // move!
}
```

std::swap

```
template <typename T>
void swap(T& a, T& b) noexcept {
    T c(std::move(a)); // move constructor
    a = std::move(b);   // move assignment
    b = std::move(c);   // move assignment
}
```